

stimpack: a modular framework for precise multisensory stimulus generation in systems neuroscience

Minseung Choi¹✉, Joshua B. Melander¹, Steven G. Herbst², Carl F. R. Wienecke^{1,3}, Heberto Mayorquin⁴, Alex Y. Hao¹, Manze Zhang¹, Jacob C. Simon¹, David D. Au¹, Bella E. Brezovec¹, Stephen A. Baccus¹, Thomas R. Clandinin^{1,5*}, and Maxwell H. Turner^{1,6*}✉

¹ Department of Neurobiology, Stanford University, USA ² Department of Electrical Engineering, Stanford University, USA ³ Department of Neurobiology, Harvard Medical School, USA ⁴ CatalystNeuro, USA ⁵ Biohub, San Francisco, USA ⁶ Department of Biological Sciences, University at Albany, SUNY, USA ✉ Corresponding author * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ✉
- [Repository](#) ✉
- [Archive](#) ✉

Editor: ✉

Submitted: 15 August 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).

Summary

A central goal of neuroscience is to understand how animals engage with the natural world. Studying this experimentally requires the ability to present stimuli that capture the complexity of dynamic sensory inputs during animal behavior. But performing such experiments across labs, and interpreting the resulting data, can be challenging because doing so requires highly flexible and precisely defined stimuli that can be synchronized with other measurement tools, and can be reconstructed to facilitate data analysis.

To address these challenges, we introduce stimpack, a Python framework for open- and closed-loop stimulus delivery. stimpack is a hardware-agnostic stimulus rendering environment where a single protocol drives perspective-corrected displays across any number and arrangement of screens, sound playback, analog outputs (e.g. optogenetic stimulation, odor delivery, or reward), and closed-loop position updates. Stimuli and devices are **modules** addressed by name over a remote-procedure-call layer, and one protocol call can invoke several modules at once. Data are written through pluggable backends, including the neuroscience community standard Neurodata Without Borders (NWB) (Rübel et al., 2022; Teeters et al., 2015).

Everything specific to a laboratory or experimental setup (e.g. screen geometry, hardware drivers, stimulus definitions) is coded in a lab- or user-specific labpack outside the installed package, discovered at runtime from a configuration file. One protocol library runs across rigs with different screen geometries and hardware, allowing experiments to be replicated precisely across laboratories.

Statement of need

Existing software packages for parametric psychophysics make naturalistic, behavior-coupled stimuli difficult to implement, while packages designed for three-dimensional virtual environments make precise parameterization challenging. stimpack was designed for laboratories that position visual displays around a head-fixed or tethered animal and couple the outputs of these displays to animal behavior in closed loop. In this context, stimpack satisfies three important requirements. First, with stimpack the same visual stimulus can be shown across different presentation geometries, including setups with multiple displays or curved screens. A stimulus can be specified in terms of visual angles relative to the subject,

as in parametric psychophysics, or specified by Euclidean positions and sizes, as in a virtual environment. In either case, each display renders the stimulus at perspective-corrected angles (Figure 2) such that quantities measured against the stimulus, from receptive fields to speed tuning, emerge in the correct angular coordinates on any rig. Second, the same package can coordinate multisensory stimulus delivery, closed-loop control, and analog device control, allowing these devices to be synchronized precisely. Finally, stimpack's files are valid NWB, so acquired data (e.g. physiological recordings, functional imaging, behavior tracking) can be added with pynwb or NeuroConv (Mayorquin et al., 2025), browsed with Neurosift (Magland et al., 2024), and archived on DANDI.

State of the field

We created stimpack to fill a gap by enabling flexible display geometries, native closed-loop control, modular device handling, and simple yet powerful stimulus design. Existing software packages for experiment control and visual stimulus delivery typically satisfy some, but not all, of these requirements.

Some visual display packages were built for **precise parametric control of a stimulus shown to a participant**. These packages make it easy to design visual stimuli, but can be difficult to extend to new screen geometries or incorporate closed-loop control based on behavioral feedback. For example, Psychtoolbox (Brainard, 1997; Pelli, 1997) and PsychoPy (Peirce et al., 2019) are designed around a display viewed roughly head-on, not around several surfaces at arbitrary orientations, and cannot readily be coupled to movement. A similar static-observer display model underlies tools for head-fixed physiology such as MonkeyLogic, PLDAPS, MWorks, and Rigbox (Bhagat et al., 2020; Eastman & Huk, 2012; Hwang et al., 2019; The MWorks Project, 2026), as well as in the Python packages Vision Egg (Straw, 2008) and QDSpy (Franke et al., 2019).

Other software packages have been purpose-built for use cases that involve **closed-loop control in a virtual environment**. These packages, while often delivering good performance, are often designed with specific hardware setups or requirements. For example, Stytra (Štih et al., 2019) and vxPy (Hladnik, 2026) have been used for a single preparation (larval zebrafish) while FreemoVR (Stowers et al., 2017) extended similar functionality to freely moving animals. Modular, custom LED arenas designed for work in fruit flies (Isaacson et al., 2022; Reiser & Dickinson, 2008) reach kilohertz binary frame rates with specific hardware requirements and a fixed geometry and resolution. Each of these packages allows for closed-loop control of parameterized stimuli, but does not directly support transferring experimental specifications to other rig geometries.

The existing packages that are most similar to stimpack can be extended to **diverse display types and geometries**. For example, ViRMEn (Aronov & Tank, 2014) and BonVision (Lopes et al., 2021) allow for curved displays, including conical and spherical screens. Bonsai (Lopes et al., 2015) handles closed-loop control and mixed devices, but trial structure and logging are left to the workflow's author, and rig calibration is embedded in the workflow itself. However, a stimpack protocol is a diff-able Python class whose parameters sweep across trials and can be stored in NWB. Finally, game engines like Unity or Unreal that have been used in neuroscience supply worlds rather than calibrated parametric stimuli, whether used directly (Jangraw et al., 2014) or wrapped for dome projection (Shapcott et al., 2025).

stimpack satisfies all these requirements, with a focus on a flexible architecture. Everything rig-specific is configured in a laboratory-specific labpack rather than inside experiment code, so **one protocol can be used across rigs and preparations**. To our knowledge, no existing package combines observer-centered perspective correction across an arbitrary number of flat screens and curved surfaces in one rig, closed-loop coupling to any locomotor tracker, non-visual outputs (sound, analog control) addressed from the same protocol, and trial-parameterized metadata saved in a community data standard by default.

Software design

One module per capability

In *stimpack*, an experiment runs as a **client**, which executes the protocol and writes metadata to the data file, and a **server**, which controls hardware devices, including display screens, as a set of modules. One command from the client can control several modules at once (Figure 1). Four modules, defined by the capability rather than a particular device, are included in *stimpack*: **visual**, which controls visual stimulus rendering and display; **audio**, which plays sound stimuli through a sound card; **locomotion**, which handles locomotion-related behavior data for closed-loop feedback; and **voltage_out**, which can send voltage waveforms to control devices (e.g. optogenetics, odor delivery, reward). This modular design choice allows a shared protocol library to run on rigs whose hardware may differ.




Figure 1: *stimpack*'s architecture. A protocol names the module for each call, and the server routes it. Inputs update a subject state that outputs follow, such that the closed loop does not pass through the client. Stacked cards are extension points: a new capability is a new module, not a change to the core.

Visual module

Displays are defined in the server code, which includes a description of the screen geometry relative to the observer. Built-in stimuli range from parametric primitives to natural movies. Stimulus shapes can be painted with monochrome or RGB textures for stimuli like gratings, spatiotemporal white noise, or movie frames. Layering, masking, and aperturing of stimuli can be achieved by a combination of alpha blending and rendering surfaces with depth.

The visual module supports two display geometries that have separate rendering paths. **Flat screens** are described by their corner locations (in meters), and each stimulus is rendered per subscreen (one perspective view per flat surface) through a generalized off-axis perspective projection (Kooima, 2009), so an object subtends the angle it should from where the subject sits. **Curved screens** (e.g. a hemisphere or a cylinder) cannot be described by a flat projection frustum, so rendering happens in two steps. The scene is first rendered into a cube centered on the subject, capturing what the subject should see in every direction. The screen's mesh is drawn in display device (e.g. projector) coordinates, and each output pixel samples the cube along the direction of the screen point it shows (Figure 2).

The visual module supports most computer display devices. Temporal multiplexing packs up to three stimulus timepoints into the color channels of each output frame, so a 120 Hz video signal can drive a 360 Hz monochrome display, approaching the temporal regime of specialized LED hardware (Reiser & Dickinson, 2008). During stimulus presentation, a small square in the corner of the display inverts on every rendered frame to enable accurate synchronization with other recorded data and to enable detection of dropped frames (Figure 2d).




Figure 2: One stimulus specification, three display geometries, one visual experience. (a) Three rigs, to scale: one monitor; two monitors meeting ahead of the subject; a hemispherical screen lit by a projector. Red: the photodiode at each display's synchronization square. (b) The frames stimpack sends to the displays for one specified scene – a 10° checkerboard left of azimuth zero, and a 360° photograph right of it: stimpack warps the images according to the specification of each display. (c) The subject's visual field: the color seen in each direction is read off the delivered frame, at the display point visible in that direction. Grey marks directions with no display surface; faint outlines trace the coverages of the other two rigs for comparison across rows. (d) Frame delivery at the photodiode: nominal inversions every 8.3 ms (120 Hz), and two dropped frames, one held bright and one held dark.

123 Audio, locomotion, and voltage output modules

124 The audio module drives a sound card. A protocol addresses the sound card in the same way
125 that the screens are addressed, naming a sound class whose waveform is rendered when the trial
126 loads and played when the trial starts. For locomotion-based closed-loop control, lab-defined
127 protocol code supplies a control function, which the server calls on every tracker update with
128 the full subject state (position, orientation, and any laboratory-defined fields measured through
129 labpack-defined trackers). The same interface lets a trial end on what the animal did, with
130 the reason recorded. The voltage-output module supplies the dispatch; the calls themselves
131 (steps, waveform streams, acquisition triggers) are defined and implemented by the labpack
132 DAQ driver.

133 Protocols, data, and the interface

134 A protocol is a Python class in the labpack that declares the parameters of a run and of the
135 stimulus. Named presets can be saved and loaded to standardize protocol parameter sets.
136 Users can also run an ensemble, which is a saved queue of protocols. Data are written through
137 pluggable backends: HDF5 by subject, series, and trial, with parameters at each level, or NWB
138 (Rübel et al., 2022; Teeters et al., 2015), the community standard. stimpack includes a GUI
139 that is a shell over these objects and identical on every rig (Figure 3).

the protocol: a Python class in the labpack (abridged)

```
class AudiovisualPairing(BaseProtocol):
    """A pulse song and a moving patch, started together."""

    def get_protocol_parameter_defaults(self):
        return {'pre_time': 0.5,
                'stim_time': 2.0,
                'tail_time': 1.0,
                'freq': [225.0, 120.0],
                'intensity': 0.0,
                'center': (0, 0),
                'speed': 60.0,
                'angle': 0.0,
                'volume': 0.5}

    def get_trial_parameters(self):
        ... # each trial's MovingPatch + PulseSong

# each trial, stimpack turns the descriptors into module calls:
server.target('visual').load_stim(name='MovingPatch', theta=..., ...)
server.target('audio').load_stim(name='PulseSong', freq=225.0, ...)
sleep(pre_time) # gray screen, silence
server.target('all').start_stim() # one start, both modalities
sleep(stim_time) # the patch sweeps, the song plays
server.target('all').stop_stim()
sleep(tail_time)
```

the GUI's Main tab, built from those declarations, mid-run

The screenshot shows the GUI's Main tab for the 'AudiovisualPairing' protocol. It displays various parameters such as 'num_trials', 'idle_color', 'all_combinations', 'randomize_order', 'pre_time', 'stim_time', 'tail_time', 'freq', 'intensity', 'center', and 'speed'. The 'freq' parameter is set to [225.0, 120.0] and is highlighted with a solid red box. Below this, the 'This trial' section shows the current frequency of 225.0, which is also highlighted with a dashed red box. The GUI includes tabs for 'Main', 'Ensemble', 'Subject', and 'File', and buttons for 'Save preset', 'Delete preset', 'View', 'Record', 'Pause', 'Stop', and 'Note'.

Figure 3: From protocol code to a running experiment. Left: the AudiovisualPairing protocol, abridged, and the module calls that stimpack makes from its descriptors on each trial – one load_stim per descriptor, routed by target; one start for both modalities; the timing declarations pacing the trial. Right: the GUI's Main tab mid-run, its parameter fields built from the class's own declarations. The list-valued freq (solid red box) sweeps across trials in randomized order, and "This trial" shows the current draw (dashed red boxes), here reaching the audio module.

140 Extensibility

141 stimpack is extensible at every layer. Stimulus classes load from the labpack at runtime
142 and are addressed exactly as built-ins; data backends, trackers, and DAQ drivers are likewise
143 labpack code. A new capability (e.g. a novel sensor or effector) is a three-line subclass of the
144 same module base class the built-ins inherit, registered with the server.

145 Research impact statement

146 stimpack consolidates four earlier Clandinin lab packages: flystim (perspective-corrected
147 rendering), flyrpc (client-server messaging), visprotocol (protocols, metadata, GUI), and
148 multistim (auditory stimulation). None has been described in an archival publication, and
149 every author of those packages is an author here. The subscreen geometry is based on that
150 of flystim. The curved-screen path was checked against flymax, an earlier MATLAB-based
151 hemisphere-projector display in the lab, and its measured photometry. Predecessor versions
152 underlie published work on fly visual processing and behavior (Currier & Clandinin, 2025; Mano
153 et al., 2023; Turner et al., 2022), the most recent of which names flystim, visprotocol, and
154 stimpack itself. In addition, stimpack has been used to reconstruct spatiotemporal receptive
155 fields in mouse retina and visual cortex (Au et al., 2026; Weddington et al., 2026), including

the example in Figure 4. The consolidation brought the modular architecture described above, improvements throughout each module, and three specific shifts. First, the viewer became a **subject with state**, which can be sent to every module, so one protocol can be used for a fly on a ball or a mouse on a treadmill. Second, **closed-loop control moved into the render loop**. Third, **laboratory-specific code left the core package**. Whereas visprotocol included lab-specific protocols and hardware drivers in its core, stimpack contains a generic core, and lab-specific code is housed in a labpack. stimpack is in use in ongoing *Drosophila* experiments in the Clandinin (Stanford), Turner (Albany), and Murthy (Princeton) laboratories, and in mouse work in the Baccus and Mitra laboratories (Stanford).




Figure 4: Spatiotemporal receptive field (STRF) of a mouse primary visual cortex neuron. The STRF of a cortical neuron recorded with a Neuropixels array was mapped with a white noise stimulus generated and rendered on two flat screens (similar to Figure 2, second row) with stimpack. The white noise stimulus had a 13° spatial correlation and a 33-ms temporal correlation. The eye was tracked to reconstruct the retinal image at a spatial scale of $< 1^\circ$ (Au et al., 2026).

AI usage disclosure

All primary architectural design decisions and implementations were made by authors. Generative AI (Claude Opus 4.8, Opus 5, and Fable 5 from Anthropic) was used for parts of code generation and editing (multisample anti-aliasing, curved screen, audio module, experiment GUI elements), code merging (curved screen, NWB data), documentation, testing scaffolding, and manuscript revision. Authors thoroughly reviewed, modified, and validated all AI-generated content.

Acknowledgements

We thank the Clandinin, Turner, Baccus, Mitra, and Murthy Labs for testing across rigs and preparations. This work was supported by Stanford Graduate Fellowships (MC, SGH), a DoD NDSEG Fellowship (MC), a Hertz Fellowship (SGH), and NIH grants R00-EY032549, R01-EY022638, R01-EY022933, R01-EY025087, R01-EY034566 and P30-EY026877. TRC is a Biohub, San Francisco, Investigator.

References

- Aronov, D., & Tank, D. W. (2014). Engagement of neural circuits underlying 2D spatial navigation in a rodent virtual reality system. *Neuron*, 84(2), 442–456. <https://doi.org/10.1016/j.neuron.2014.08.042>
- Au, D. D., Melander, J. B., Weddington, J. C., Faragalla, Y., Alaoui, Z., Liu, S., Xu, Q., & Baccus, S. A. (2026). Single-camera, calibration-free gaze estimation using corneal

- 184 reflections. *bioRxiv*. <https://doi.org/10.64898/2026.06.08.731021>
- 185 Bhagat, J., Wells, M. J., Harris, K. D., Carandini, M., & Burgess, C. P. (2020).
186 Rigbox: An open-source toolbox for probing neurons and behavior. *eNeuro*, 7(4),
187 ENEURO.0406–19.2020. <https://doi.org/10.1523/ENEURO.0406-19.2020>
- 188 Brainard, D. H. (1997). The Psychophysics Toolbox. *Spatial Vision*, 10(4), 433–436. <https://doi.org/10.1163/156856897X00357>
- 190 Currier, T. A., & Clandinin, T. R. (2025). Infrequent strong connections constrain connectomic
191 predictions of neuronal function. *Cell*, 188(16), 4366–4381.e14. <https://doi.org/10.1016/j.cell.2025.05.007>
- 193 Eastman, K. M., & Huk, A. C. (2012). PLDAPS: A hardware architecture and software toolbox
194 for neurophysiology requiring complex visual stimuli and online behavioral control. *Frontiers*
195 *in Neuroinformatics*, 6, 1. <https://doi.org/10.3389/fninf.2012.00001>
- 196 Franke, K., Maia Chagas, A., Zhao, Z., Zimmermann, M. J. Y., Bartel, P., Qiu, Y., Szatko,
197 K. P., Baden, T., & Euler, T. (2019). An arbitrary-spectrum spatial visual stimulator for
198 vision research. *eLife*, 8, e48779. <https://doi.org/10.7554/eLife.48779>
- 199 Hladnik, T. (2026). *vxPy: Multiprocess-based software for vision experiments in Python*.
200 <https://github.com/thladnik/vxPy>
- 201 Hwang, J., Mitz, A. R., & Murray, E. A. (2019). NIMH MonkeyLogic: Behavioral control
202 and data acquisition in MATLAB. *Journal of Neuroscience Methods*, 323, 13–21. <https://doi.org/10.1016/j.jneumeth.2019.05.002>
- 204 Isaacson, M., Ferguson, L., Loesche, F., Ganguly, I., Chen, J., Chiu, A., Liu, J., Dickson,
205 W., & Reiser, M. (2022). A high-speed, modular display system for diverse neuroscience
206 applications. *bioRxiv*. <https://doi.org/10.1101/2022.08.02.502550>
- 207 Jangraw, D. C., Johri, A., Gribetz, M., & Sajda, P. (2014). NEDE: An open-source scripting
208 suite for developing experiments in 3D virtual environments. *Journal of Neuroscience*
209 *Methods*, 235, 245–251. <https://doi.org/10.1016/j.jneumeth.2014.06.033>
- 210 Kooima, R. (2009). *Generalized perspective projection*. Louisiana State University.
- 211 Lopes, G., Bonacchi, N., Frazão, J., Neto, J. P., Atallah, B. V., Soares, S., Moreira, L., Matias,
212 S., Itskov, P. M., Correia, P. A., Medina, R. E., Calcaterra, L., Dreosti, E., Paton, J. J., &
213 Kampff, A. R. (2015). Bonsai: An event-based framework for processing and controlling data
214 streams. *Frontiers in Neuroinformatics*, 9, 7. <https://doi.org/10.3389/fninf.2015.00007>
- 215 Lopes, G., Farrell, K., Horrocks, E. A. B., Lee, C.-Y., Morimoto, M. M., Muzzu, T.,
216 Papanikolaou, A., Rodrigues, F. R., Wheatcroft, T., Zucca, S., Solomon, S. G., & Saleem,
217 A. B. (2021). Creating and controlling visual environments using BonVision. *eLife*, 10,
218 e65541. <https://doi.org/10.7554/eLife.65541>
- 219 Magland, J., Soules, J., Baker, C., & Dichter, B. (2024). Neurosift: DANDI exploration
220 and NWB visualization in the browser. *Journal of Open Source Software*, 9(97), 6590.
221 <https://doi.org/10.21105/joss.06590>
- 222 Mano, O., Choi, M., Tanaka, R., Creamer, M. S., Matos, N. C. B., Shomar, J. W., Badwan,
223 B. A., Clandinin, T. R., & Clark, D. A. (2023). Long-timescale anti-directional rotation in
224 *Drosophila* optomotor behavior. *eLife*, 12, e86076. <https://doi.org/10.7554/eLife.86076>
- 225 Mayorquin, H., Baker, C., Adkisson-Floro, P., Weigl, S., Trapani, A., Tauffer, L., Rübél,
226 O., & Dichter, B. (2025). NeuroConv: Streamlining neurophysiology data conversion
227 to the NWB standard. *Proceedings of the 24th Python in Science Conference*. <https://doi.org/10.25080/cehj4257>
- 229 Peirce, J., Gray, J. R., Simpson, S., MacAskill, M., Höchenberger, R., Sogo, H., Kastman,
230 E., & Lindeløv, J. K. (2019). PsychoPy2: Experiments in behavior made easy. *Behavior*

- 231 *Research Methods*, 51(1), 195–203. <https://doi.org/10.3758/s13428-018-01193-y>
- 232 Pelli, D. G. (1997). The VideoToolbox software for visual psychophysics: Transforming numbers
233 into movies. *Spatial Vision*, 10(4), 437–442. <https://doi.org/10.1163/156856897X00366>
- 234 Reiser, M. B., & Dickinson, M. H. (2008). A modular display system for insect behavioral
235 neuroscience. *Journal of Neuroscience Methods*, 167(2), 127–139. <https://doi.org/10.1016/j.jneumeth.2007.07.019>
- 236
- 237 Rübél, O., Tritt, A., Ly, R., Dichter, B. K., Ghosh, S., Niu, L., Baker, P., Soltesz, I., Ng,
238 L., Svoboda, K., Frank, L., & Bouchard, K. E. (2022). The Neurodata Without Borders
239 ecosystem for neurophysiological data science. *eLife*, 11, e78362. <https://doi.org/10.7554/eLife.78362>
- 240
- 241 Shapcott, K. A., Weigand, M., Glukhova, M., Havenith, M. N., & Schölvinck, M. L. (2025).
242 DomeVR: Immersive virtual reality for primates and rodents. *PLOS ONE*, 20(1), e0308848.
243 <https://doi.org/10.1371/journal.pone.0308848>
- 244 Štíh, V., Petrucco, L., Kist, A. M., & Portugues, R. (2019). Stytra: An open-source,
245 integrated system for stimulation, tracking and closed-loop behavioral experiments. *PLOS*
246 *Computational Biology*, 15(4), e1006699. <https://doi.org/10.1371/journal.pcbi.1006699>
- 247 Stowers, J. R., Hofbauer, M., Bastien, R., Griessner, J., Higgins, P., Farooqui, S., Fischer,
248 R. M., Nowikovsky, K., Haubensak, W., Couzin, I. D., Tessmar-Raible, K., & Straw, A.
249 D. (2017). Virtual reality for freely moving animals. *Nature Methods*, 14(10), 995–1002.
250 <https://doi.org/10.1038/nmeth.4399>
- 251 Straw, A. D. (2008). Vision Egg: An open-source library for realtime visual stimulus generation.
252 *Frontiers in Neuroinformatics*, 2, 4. <https://doi.org/10.3389/neuro.11.004.2008>
- 253 Teeters, J. L., Godfrey, K., Young, R., Dang, C., Friedsam, C., Wark, B., Asari, H., Peron,
254 S., Li, N., Peyrache, A., Denisov, G., Siegle, J. H., Olsen, S. R., Martin, C., Chun, M.,
255 Tripathy, S., Blanche, T. J., Harris, K., Buzsáki, G., ... Sommer, F. T. (2015). Neurodata
256 without borders: Creating a common data format for neurophysiology. *Neuron*, 88(4),
257 629–634. <https://doi.org/10.1016/j.neuron.2015.10.025>
- 258 The MWorks Project. (2026). *MWorks: A framework for conducting realtime neurophysiology*
259 *and psychophysics experiments*. <https://mworks.github.io/>
- 260 Turner, M. H., Krieger, A., Pang, M. M., & Clandinin, T. R. (2022). Visual and motor
261 signatures of locomotion dynamically shape a population code for feature detection in
262 *Drosophila*. *eLife*, 11, e82587. <https://doi.org/10.7554/eLife.82587>
- 263 Weddington, J. C., Au, D. D., Melander, J. B., Faragalla, Y., Alaoui, Z., & Baccus, S. A.
264 (2026). Parallel construction of object motion from retina to cortex. *bioRxiv*. <https://doi.org/10.64898/2026.07.27.740872>
- 265