

PySlyde: A Lightweight, Open-Source Toolkit for Pathology Preprocessing

Gregory Verghese^{1,2*¶}, Anthony Baptista^{2,3*}, Chima Eke^{2*}, Holly Rafique^{2*}, Elizabeth Ing-Simmons^{1,4*}, Enrico Parisini^{1,2}, Mengyuan Li^{1,2}, Fathima Mohamed², Ananya Bhalla^{2,5}, Lucy Ryan², Michael Pitcher^{1,2}, Concetta Piazzese^{1,6}, James Graham^{1,4}, Dinis Pedro Calado⁵, Christopher R. S. Banerji^{2,3}, and Anita Grigoriadis^{1,2*¶}

¹ PharosAI, London, UK ² Cancer Bioinformatics, School of Cancer and Pharmaceutical Sciences, Faculty of Life Sciences and Medicine, King's College London, London, UK ³ The Alan Turing Institute, The British Library, London, UK ⁴ eResearch, King's College London, London, UK ⁵ The Francis Crick Institute, London, UK ⁶ Barts Life Sciences, Barts Health NHS Trust, London, UK ¶ Corresponding author * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: Julia Romanowska

Reviewers:

- [@manunicholasjacob](#)
- [@showdawn](#)
- [@moerlemans](#)

Submitted: 09 March 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/))

Summary

Artificial intelligence (AI) represents a transformative opportunity for precision medicine, with the potential to radically improve patient outcomes (Song et al., 2023; Verghese, Lennerz, et al., 2023). Advances in computational pathology are enabling the analysis of complex histopathological data at unprecedented scale and depth. By combining modern computer vision, machine learning, and pathology, researchers can predict quantitative, molecular (Dey et al., 2025) and digital biomarkers, detect morphological and tissue patterns (S. Chen et al., 2026), and develop predictive models for diagnostics, prognosis (Verghese, Li, et al., 2023) and treatment response to complement human expertise (Laak et al., 2021).

Whole slide images (WSIs), resulting from the digitisation of histopathology slides, offer valuable spatial and morphological insights that can deepen our understanding of cancer biology. However, their size (gigapixel-scale), complexity, and variability introduce significant computational and standardisation challenges for downstream tasks (McGenity et al., 2023). Preprocessing WSIs, including tissue detection, tessellation (splitting into smaller tiled images), stain normalisation, and annotation parsing, are essential steps before any downstream AI analysis can occur (Pocock et al., 2022). Building robust and reproducible pipelines for WSI preprocessing is fundamental to computational pathology research. However, existing workflows often rely on ad hoc proprietary scripts, fragmented tools, and inconsistent formats, making standardisation and reproducibility difficult.

PySlyde is a lightweight, open-source Python toolkit built on top of OpenSlide (Goode et al., 2013), designed to address this gap and provide an intuitive approach to quickly preprocess WSIs. It provides a simple API to perform WSI loading, annotation handling, tissue detection, tile generation, and feature extraction with support for the latest histopathology foundation models. PySlyde aims to lower the barrier to entry for pathology researchers, standardise preprocessing workflows, and accelerate the development of AI-ready datasets for computational pathology research, enabling researchers to focus on model development and other downstream tasks.




Figure 1: Schematic overview of the WSI processing workflow supported by PySlyde, illustrating the core functionalities across the main classes and modules `Slide`, `WSIParser`, `FeatureGenerator`, and `IO`. Top row, from left to right: (i) WSI loading using compatible formats through `OpenSlide`; (ii) annotation mask generation with support for `QuPath` or `ImageJ` formats; and (iii) preprocessing steps such as tissue detection. Bottom row, from left to right: (iv) tessellation of WSIs into tissue tiles using the `WSIParser` class; (v) feature extraction from tiles with support for pretrained encoders such as `Virchow2`, `H-optimus`, `Gigapath`, or `UNI`; and (vi) structured storage of tile embeddings and metadata using `RocksDB`, `NumPy`, or `LMDB`.

Statement of Need

WSI preprocessing is a foundational but time-consuming step in computational pathology. Common tasks, such as annotation parsing, binary and annotation mask generation, stain normalisation, and tessellation, often involve custom scripts, multiple packages, and ad hoc pipelines built for a single dataset or project. This fragmentation hinders reproducibility, increases technical debt, and slows the development and translation of AI methods into clinical workflows. The lack of standardised preprocessing practices thus remains one of the key barriers preventing the seamless integration of AI into diagnostic pathology (Bilal et al., 2025).

While powerful frameworks exist for model training and inference, few provide a simple, consistent interface for preparing WSIs at scale. PySlyde directly addresses this unmet need by providing a modular, well-documented, and extensible preprocessing toolkit that integrates seamlessly with existing ecosystems, including `OpenSlide`, `NumPy`, `PyTorch`, and modern foundation models, enabling researchers to build reproducible, scalable pipelines with minimal code.

State of the Field

Several open-source tools support whole-slide image analysis in computational pathology. At the foundational level, `OpenSlide` (Goode et al., 2013) provides a widely adopted C library and Python bindings for programmatic access to WSI formats. It enables efficient tile-based image loading, thumbnail generation, and extraction of slide metadata across multiple vendor-specific file formats. While it provides essential low-level functionality for accessing WSI data, higher-level workflows such as annotation management, spatial analysis, and feature extraction must typically be implemented by downstream libraries.

Building upon such low-level infrastructure, frameworks such as `TIAToolbox` (Pocock et al., 2022) provide comprehensive ecosystems for tissue analytics, model development, and

64 deployment. These frameworks offer extensive end-to-end functionality including model training,
65 inference, and visualisation. In contrast, PySlyde is designed specifically as a lightweight,
66 modular preprocessing layer for WSI analysis, focusing on efficient data preparation, annotation
67 handling, and feature embedding generation for downstream machine learning workflows.

68 PySlyde provides native support for annotations generated by widely adopted WSI analysis tools
69 such as QuPath (Bankhead et al., 2017) and ImageJ (Schneider et al., 2012). Annotation files
70 (e.g., QuPath JSON exports) are parsed into structured spatial representations that preserve
71 region geometry and associated metadata. These annotations can be converted into binary or
72 multi-class masks for pixel-level preprocessing, or represented as GeoDataFrame polygon objects
73 that retain slide-level spatial coordinates. Representing annotations in this format enables
74 interoperability with geospatial libraries such as GeoPandas and Shapely, facilitating spatial
75 analysis workflows that extend beyond traditional image processing pipelines. Annotations
76 are exported to interoperable formats such as GeoJSON, enabling integration with external
77 visualisation tools and GIS-based analytical workflows.

78 PySlyde further differentiates itself through its native support for histopathology foundation
79 models and structured feature extraction. Embedding extraction from modern histopathology
80 foundation models—Virchow2 (Zimmermann et al., 2024), H-optimus (Saillard et al., 2024),
81 UNI (R. J. Chen et al., 2024), and cTransPath (Wang et al., 2022)—is treated as a first-class
82 preprocessing component rather than an auxiliary utility. Extracted embeddings and associated
83 metadata are exported to scalable key-value storage backends such as LMDB and RocksDB,
84 enabling efficient indexed retrieval for large-scale machine learning and representation learning
85 workflows.

86 PySlyde addresses a specific gap in the current ecosystem: providing a minimal, modular
87 preprocessing backbone that integrates WSI annotation handling, geospatial representations,
88 and foundation model feature extraction within a single lightweight framework. This design
89 enables researchers to incorporate modern foundation model embeddings and spatially aware
90 annotation data into scalable machine learning pipelines without adopting a full end-to-end
91 pathology framework.

92 Software Design

93 PySlyde is designed as a lightweight and modular preprocessing framework for WSI analysis.
94 The architecture emphasises separation of functionality between slide representation, spatial
95 parsing, feature extraction, and data persistence. This design allows individual components of
96 the preprocessing pipeline to be used independently while maintaining a coherent workflow
97 for large-scale computational pathology experiments. Rather than providing a tightly coupled
98 end-to-end pipeline, PySlyde prioritises modularity, enabling researchers to integrate specific
99 preprocessing stages into existing machine learning workflows.

100 At the core of the architecture is the `Slide` abstraction, which represents a single WSI together
101 with its associated spatial metadata, annotations, and derived data. The `Slide` class inherits
102 directly from the `OpenSlide` Python interface, allowing PySlyde to extend low-level slide access
103 with higher-level functionality for annotation management, spatial masking, and tissue-aware
104 preprocessing. This approach avoids re-implementing foundational image I/O functionality
105 while preserving compatibility with the broad range of vendor-specific formats supported by
106 `OpenSlide`.

107 The remaining architecture is organised around three complementary components. The
108 `WSIParser` class handles spatial partitioning and tile extraction, separating tiling logic from
109 slide representation so that large gigapixel images are processed efficiently under different
110 region-selection and magnification strategies. The `FeatureGenerator` class encapsulates
111 feature extraction workflows for modern histopathology foundation models, allowing embedding
112 generation to remain decoupled from upstream slide parsing and tiling operations. The `I/O`
113 module manages export and persistence of tiles, embeddings, and metadata through both

file-based outputs and scalable key-value stores such as Lightning Memory Mapped Database (LMDB) and RocksDB, supporting reproducible downstream analysis.

Several architectural trade-offs guided this design. First, PySlyde prioritises modularity over a monolithic pipeline structure. While tightly integrated frameworks simplify end-to-end workflows, they tend to constrain researchers to predefined processing stages. PySlyde instead provides composable abstractions that support experimentation with different tiling strategies, preprocessing methods, storage backends, and feature extractors. Second, PySlyde deliberately builds on established libraries such as OpenSlide and HistoQC rather than re-implementing low-level functionality. This reduces maintenance overhead and improves interoperability with existing pathology infrastructure, while allowing development effort to focus on annotation-aware preprocessing and foundation model integration. Together, these design choices support scalable and reproducible preprocessing workflows for computational pathology research, particularly where large WSI collections must be transformed into structured datasets for modern machine learning and representation learning applications.

Research Impact Statement

PySlyde has been used in multiple ongoing computational pathology projects within the Cancer Bioinformatics Team at King's College London, including multimodal outcome prediction studies in breast, head and neck, and pan-cancer cohorts, with several manuscripts currently under review or in preparation. It is also used within the PharosAI Research Ventures Catalyst programme for large-scale histopathology data curation across collaborating institutions including Guy's and St Thomas' NHS Foundation Trust, Queen Mary University of London, Barts Health NHS Trust, and King's College London. These applications, alongside a poster at the *ESMO AI Congress, Berlin 2025*, demonstrate its immediate research utility and support expectations for wider adoption within computational pathology workflows.

Use of Generative AI

No generative AI tools were used in the development of the original PySlyde codebase. More recently, AI-assisted development tools (e.g., Cursor v2.4.31) have been used to support limited refactoring (e.g., type annotations) and code maintenance tasks. Recent AI-assisted changes were reviewed through standard code review and CI/CD validation processes on GitHub to ensure correctness and reliability. Generative AI tools were also used to support the writing, reviewing, and editing of this manuscript, with all final content verified and edited by all the authors.

Acknowledgements

We thank members of the King's College London Cancer Bioinformatics Team, the PharosAI Research Ventures Catalyst Programme, and collaborators at The Alan Turing Institute, The Francis Crick Institute, and Barts Life Sciences for feedback, testing, and support. Holly Rafique, Pre-Doctoral Fellow NIHR303406, is funded by the NIHR. Anita Grigoriadis, Gregory Verghese, Mengyuan Li, Elizabeth Ing-Simmons, and Enrico Parisini are funded by the UK government, Research Ventures Catalyst Programme, Department of Science Innovation and Technology, and the Guy's Cancer Charity under PharosAI for this research project. Anthony Baptista, Chima Eke, and Anita Grigoriadis acknowledge support from the CRUK City of London Centre Award [CTRQQR-2021/100004]. Christopher R.S. Banerji is supported by a King's College London AI+ Fellowship. Dinis Pedro Calado acknowledges support from The Francis Crick Institute, which receives core funding from Cancer Research UK (CC2078), the UK Medical Research Council (CC2078), and the Wellcome Trust (CC2078). Dinis Pedro Calado and Anita Grigoriadis acknowledge support from the UK Medical Research Council

(MR/W025221/1). Lucy Ryan is funded by the Pathological Society of Great Britain and Ireland and the Jean Shanks Foundation (JSPS CPF 1023 05). The views expressed in this publication are those of the author(s) and not necessarily those of the NIHR.

References

- Bankhead, P., Loughrey, M. B., Fernández, J. A., Dombrowski, Y., McArt, D. G., Dunne, P. D. J., McQuaid, S., Gray, R. T., Murray, L. J., Coleman, H. G., James, J. A., Salto-Tellez, M., & Hamilton, P. (2017). QuPath: Open source software for digital pathology image analysis. *Scientific Reports*, 7. <https://api.semanticscholar.org/CorpusID:12778254>
- Bilal, M., Aadam, Raza, M., Altherwy, Y., Alsuhaibani, A., Abduljabbar, A., Almarshad, F., Golding, P., & Rajpoot, N. (2025). *Foundation models in computational pathology: A review of challenges, opportunities, and impact*. <https://arxiv.org/abs/2502.08333>
- Chen, R. J., Ding, T., Lu, M. Y., Williamson, D. F. K., Jaume, G., Song, A. H., Chen, B., Zhang, A., Shao, D., Shaban, M., Williams, M., Oldenburg, L., Weishaupt, L. L., Wang, J. J., Vaidya, A., Le, L. P., Gerber, G. K., Sahai, S., Williams, W., & Mahmood, F. (2024). Towards a general-purpose foundation model for computational pathology. *Nature Medicine*, 30, 850–862. <https://api.semanticscholar.org/CorpusID:268534221>
- Chen, S., Parreno-Centeno, M., Booker, G., Verghese, G. E., Mohamed, F. S., Arslan, S., Pandya, P., Oozeer, A., D'Angelo, M., Barrow, R., Nelán, R., Sobral-Leite, M., Martino, F. D., Briskén, C., Smalley, M. J., Lips, E. H., Gillett, C., Jones, L. J., Banerji, C. R. S., ... Grigoriadis, A. (2026). Normal breast tissue (NBT)-classifiers: Advancing compartment classification in normal breast histology. *bioRxiv*. <https://api.semanticscholar.org/CorpusID:278285427>
- Dey, S., Banerji, C. R. S., Basuchowdhuri, P., Saha, S. K., Parashar, D., & Chakraborti, T. (2025). Generating crossmodal gene expression from cancer histopathology improves multimodal AI predictions. *Nature Communications*, 17. <https://api.semanticscholar.org/CorpusID:276095062>
- Goode, A., Gilbert, B., Harkes, J., Jukic, D., & Satyanarayanan, M. (2013). OpenSlide: A vendor-neutral software foundation for digital pathology. *Journal of Pathology Informatics*, 4(1), 27. <https://doi.org/10.4103/2153-3539.119005>
- Laak, J. A. van der, Litjens, G. J. S., & Ciompi, F. (2021). Deep learning in histopathology: The path to the clinic. *Nature Medicine*, 27, 775–784. <https://api.semanticscholar.org/CorpusID:234597294>
- McGenity, C., Clarke, E. L., Jennings, C., Matthews, G., Cartlidge, C., Fredueh-Agyemang, H., Stocken, D. D., & Treanor, D. (2023). Artificial intelligence in digital pathology: A systematic review and meta-analysis of diagnostic test accuracy. *NPJ Digital Medicine*, 7. <https://api.semanticscholar.org/CorpusID:269587404>
- Pocock, J., Graham, S., Vu, Q. D., Jahanifar, M., Deshpande, S., Hadjigeorgiou, G., Shephard, A. J., Bashir, R. M. S., Bilal, M., Lu, W., Epstein, D. B. A., Minhas, F. A., Rajpoot, N. M., & Raza, S. E. A. (2022). TIAToolbox as an end-to-end library for advanced tissue image analytics. *Communications Medicine*, 2. <https://api.semanticscholar.org/CorpusID:252498570>
- Saillard, C., Jenatton, R., Llinares-López, F., Mariet, Z., Cahané, D., Durand, E., & Vert, J.-P. (2024). *H-optimus-0*. <https://github.com/bioptimus/releases/tree/main/models/h-optimus/v0>
- Schneider, C. A., Rasband, W. S., & Eliceiri, K. W. (2012). NIH image to ImageJ: 25 years of image analysis. *Nature Methods*, 9, 671–675. <https://api.semanticscholar.org/CorpusID:13131970>

- 207 Song, A. H., Jaume, G., Williamson, D. F. K., Lu, M. Y., Vaidya, A., Miller, T. R., & Mahmood,
208 F. (2023). Artificial intelligence for digital and computational pathology. *Nature Reviews*
209 *Bioengineering*, 1, 930–949. <https://api.semanticscholar.org/CorpusID:263651769>
- 210 Verghese, G. E., Lennerz, J. K., Ruta, D., Ng, W., Thavaraj, S., Siziopikou, K. P., Naidoo, T.,
211 Rane, S. U., Salgado, R., Pinder, S. E., & Grigoriadis, A. (2023). Computational pathology
212 in cancer diagnosis, prognosis, and prediction – present day and prospects. *The Journal of*
213 *Pathology*, 260, 551–563. <https://api.semanticscholar.org/CorpusID:260898407>
- 214 Verghese, G. E., Li, M., Liu, F., Lohan, A., Kurian, N. C., Meena, S., Gazinska, P., Shah, A.,
215 Oozeer, A., Chan, T. W. S., Opdam, M., Linn, S. C., Gillett, C., Alberts, E., Hardiman,
216 T., Jones, S., Thavaraj, S., Jones, J. L., Salgado, R., ... Grigoriadis, A. (2023). Multiscale
217 deep learning framework captures systemic immune features in lymph nodes predictive of
218 triple negative breast cancer outcome in large-scale studies. *The Journal of Pathology*, 260,
219 376–389. <https://api.semanticscholar.org/CorpusID:258908700>
- 220 Wang, X., Yang, S., Zhang, J., Wang, M., Zhang, J., Yang, W., Huang, J., & Han, X.
221 (2022). Transformer-based unsupervised contrastive learning for histopathological image
222 classification. *Medical Image Analysis*, 81, 102559. <https://api.semanticscholar.org/CorpusID:251207603>
- 224 Zimmermann, E., Vorontsov, E., Viret, J., Casson, A., Zelechowski, M., Shaikovski, G.,
225 Tenenholtz, N., Hall, J., Fuchs, T. J., Fusi, N., Liu, S., & Severson, K. (2024). Virchow2:
226 Scaling self-supervised mixed magnification models in pathology. *ArXiv*, abs/2408.00738.
227 <https://api.semanticscholar.org/CorpusID:271600873>