

torchTT: A PyTorch-based Tensor-Train Toolbox

Ion Gabriel Ion ¹

¹ Terra Quantum AG

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Matt Graham](#) 

Reviewers:

- [@dsuess](#)
- [@MarieRoald](#)

Submitted: 16 January 2026

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).

Summary

torchTT is a Python library enabling Tensor-Train (TT) decomposition (I. V. Oseledets, 2011) within the PyTorch framework (Paszke et al., 2019). It allows users to create and manipulate high-dimensional tensors in compressed TT format using PyTorch-like syntax and semantics. Operations such as basic linear algebra, matrix-vector products, and Kronecker products can be performed directly on the compressed representation without needing to form the full tensor. All computations leverage PyTorch's capabilities: TT instances live on the same device as their underlying data, facilitating seamless GPU acceleration, and automatic differentiation works through TT operations. Under the hood, torchTT exploits low-rank structure in high-dimensional arrays, dramatically reducing memory and computation compared to dense tensors. These features make torchTT especially suitable for problems with millions or billions of degrees of freedom where explicit storage is infeasible.

Statement of Need

High-dimensional tensors arise in many areas of science and engineering, such as the discretization of d -dimensional PDEs, state spaces for ODE systems, and many-body wavefunctions in quantum chemistry. Naively storing or computing with a full tensor scales exponentially in d (the curse of dimensionality). The TT decomposition combats this by factorizing a tensor into a sequence of smaller, low-rank 3D cores, reducing the storage from $\mathcal{O}(n^d)$ to $\mathcal{O}(dnr^2)$ where r is the TT rank. torchTT targets researchers and engineers in scientific computing, computational physics, and deep learning who need to work with such high-dimensional objects. It addresses the growing need for a single library that combines hardware-accelerated tensor network algebra with advanced numerical algorithms within the PyTorch ecosystem, enabling seamless integration into modern machine learning workflows.

State of the Field

Several open-source packages exist for tensor network computations and the TT format specifically. The ttpy library (I. Oseledets, 2021) is one of the earliest and most comprehensive Python implementations for TT decomposition, offering a wide array of functions including linear solvers and cross-approximation. However, ttpy relies on NumPy and SciPy, meaning it lacks native GPU acceleration and automatic differentiation, which are critical for modern deep learning workflows.

More recently, libraries like tntorch (Ballester-Ripoll, 2022), t3f (Novikov et al., 2020), and TensorLy (Kossaifi et al., 2019) have bridged the gap between tensor networks and machine learning frameworks by utilizing PyTorch and TensorFlow. These packages provide GPU acceleration and autograd capabilities. While tntorch offers cross-approximation routines, none of these libraries provide the AMEn (Alternating Minimal Energy) solver for linear systems in the TT format nor the AMEn-based adaptive cross interpolation scheme, both of which are essential for robust numerical analysis of high-dimensional problems.

41 Additionally, in the domain of quantum physics, libraries such as TeNPy (Hauschild & Pollmann,
42 2018) provide highly optimized tools for simulating many-body systems using Matrix Product
43 States (the physics equivalent of the TT format). However, these are deeply specialized for
44 physics applications, such as the density matrix renormalization group (DMRG) algorithm,
45 rather than general-purpose numerical mathematics or deep learning.

46 torchTT addresses this gap by combining the best of both worlds. It provides a native PyTorch
47 implementation with seamless GPU acceleration and automatic differentiation, making it ideal
48 for deep learning. Concurrently, it offers advanced numerical capabilities including both DMRG
49 and AMEn-based cross approximation, and the AMEn linear system solver, positioning it
50 as a uniquely complete tool for scientific machine learning. Rather than contributing these
51 features to existing packages—which would require significant architectural changes to libraries
52 not designed for advanced iterative solvers—torchTT was built from the ground up with a
53 unified TT class that natively supports both scientific computing workflows and deep learning
54 integration.

55 Software Design

56 A central design trade-off in torchTT is balancing the expressiveness needed for advanced
57 numerical algorithms with seamless integration into the PyTorch ecosystem. Existing scientific
58 TT libraries (e.g., ttpy) are built on NumPy and thus cannot leverage GPU acceleration
59 or automatic differentiation. Conversely, ML-oriented tensor libraries prioritize autograd
60 compatibility but lack the algorithmic depth needed for iterative solvers. torchTT resolves this
61 tension through two key architectural decisions.

62 Unified TT class and PyTorch-native design

63 The central component is the `torchtt.TT` class, which provides a single, unified representation
64 for both TT-tensors and TT-matrices (tensor operators). This design choice—rather than
65 separate classes—allows algorithms like AMEn and cross-approximation to operate generically
66 on any TT object. The class stores its TT cores as standard `torch.Tensor` objects, which
67 means PyTorch's autograd graph tracks all TT operations automatically. This enables end-to-
68 end gradient-based training of models that include TT-compressed layers or TT-based loss
69 functions, without requiring custom backward passes.

- 70 ■ **Python-like API:** Operations such as addition (+), subtraction (-), elementwise
71 multiplication (*), and matrix multiplication (@) are overloaded to work directly on
72 `torchtt.TT` instances, keeping user code concise.
- 73 ■ **Device management:** Standard PyTorch methods like `.to(device)`, `.cuda()`, and
74 `.cpu()` are supported, enabling TT objects and their computations to run on CPU or
75 GPU without code changes.

76 Dual backend strategy

77 For performance-critical inner loops—particularly the AMEn solver—torchTT optionally
78 delegates to a compiled C++ extension via PyTorch's custom operator mechanism. When
79 the extension is unavailable, all algorithms fall back to a pure-PyTorch implementation with
80 identical semantics. This dual-backend strategy ensures broad portability (any platform with
81 PyTorch) while offering near-native performance when the C++ extension is compiled.

82 Functionality and Features

83 Beyond the core class, the torchTT library is organized into specialized modules targeting
84 specific scientific and machine learning applications:

- 85 ■ **Linear System Solvers** (`torchtt.solvers`): Implements advanced solvers for linear
86 systems $Ax = b$ directly in the TT format, including the AMEn method (Dolgov &
87 Savostyanov, 2014). These solvers dynamically adapt the TT ranks during iteration to
88 maintain accuracy. The module also includes elementwise division and inversion routines.
- 89 ■ **Cross Approximation** (`torchtt.interpolate`): Provides adaptive TT-cross interpolation
90 methods (I. V. Oseledets, 2010). This module constructs TT approximations of
91 multivariate functions from black-box evaluations, essential for tasks where the tensor is
92 too large to be formed explicitly but can be sampled.
- 93 ■ **Neural Network Layers** (`torchtt.nn`): Defines PyTorch-compatible neural network
94 layers. `LinearLayerTT` parametrizes dense weight matrices as TT-matrices, enabling
95 massive parameter reduction. `CompressedTTLayer` extends this concept by operating
96 directly on TT-formatted inputs and applying nonlinear activations between TT cores
97 during multiplication, targeting deep TT network architectures. All layers inherit from
98 `torch.nn.Module`, making them drop-in replacements for standard layers in deep learning
99 pipelines.
- 100 ■ **Manifold Optimization** (`torchtt.manifold`): Offers tools for Riemannian optimization
101 on the manifold of tensors with fixed TT-rank. This includes projections onto the tangent
102 space and Riemannian gradient calculation, facilitating advanced optimization tasks like
103 tensor completion.

104 Example: Solving a 4D PDE

105 The following example demonstrates how to solve a 4-dimensional Poisson equation $\Delta u = f$
106 on $[0, 1]^4$ with zero boundary conditions using `torchTT`. The right-hand side f is constructed
107 using cross-approximation, and the system is solved directly in the TT format using the AMEn
108 solver. Additional runnable examples are provided in the repository's `examples/` folder.

```
import torchtt as tntt
import torch as tn

# Define the problem size
N, d = 64, 4

# Construct the 1D Laplacian operator (finite difference)
L1d = (tn.diag(tn.ones(N-1), -1) + tn.diag(tn.ones(N-1), 1) - \
        2*tn.eye(N)) / (1/(N-1))**2
L1d[0,1] = L1d[-1,-2] = 0 # Boundary conditions
L1d_tt = tntt.TT(L1d, shape=[(N,N)])

# Construct the d-dimensional Laplacian
L = L1d_tt ** tntt.eye([N]*3) + tntt.eye([N]) ** L1d_tt ** tntt.eye([N]*2) + \
    tntt.eye([N]*2) ** L1d_tt ** tntt.eye([N]) + tntt.eye([N]*3) ** L1d_tt

# Create the grid for cross approximation
x = tntt.meshgrid([tn.linspace(0,1,N)]*d)

# Approximate the RHS f(x) = x1(x1-1)*...*xd(xd-1)
f = tntt.interpolate.function_interpolate(lambda x: tn.prod(x * (x - 1), dim=1), x, eps=

# Solve the linear system Ax = f using AMEn on GPU
u = tntt.solvers.amen_solve(L.to('cuda'), f.to('cuda'), eps=1e-6)
```

Research Impact Statement

torchTT has demonstrated realized research impact across multiple domains, as evidenced by its use in peer-reviewed publications and ongoing collaborations:

- **High-Dimensional PDEs/ODEs:** torchTT can solve PDEs and dynamical systems in high dimensions by representing solution spaces and operators in TT form. Built-in solvers like AMEn make it practical to compute time steps of multi-dimensional PDEs, breaking the curse of dimensionality. Recent applications include isogeometric analysis for parameter-dependent geometries (Ion et al., 2022) and solving the chemical master equation for parameter inference (Ion & others, 2022).
- **Computational Physics and Chemistry:** The library enables efficient representation of high-dimensional objects like quantum many-body wavefunctions. Operators can be applied via TT-matrix-vector products, and observables computed via TT inner products, facilitating large-scale simulations. Furthermore, it supports analyzing entanglement scaling in Matrix Product State (MPS) representations and constructing shallow quantum circuits (Bohun et al., 2024).
- **Machine Learning with Compressed Models:** TT-layers enable significant compression of model parameters in deep learning. torchTT's compatibility with autograd allows for end-to-end training of compressed architectures and efficient on-device inference.
- **Data Science and Function Approximation:** The library is useful for surrogate modeling and uncertainty quantification, where high-dimensional response surfaces can be approximated adaptively using cross-approximation routines. Novel methods for local interpolation allow for constructing fine-scale TT representations from coarse-grid approximations with error guarantees (Guzman & Tiunov, 2026). This includes quantitative finance applications, where TT-cross approximation provides efficient option pricing surrogates for high-dimensional risk management (Gribben et al., 2026).

Acknowledgements

The authors thank the Terra Quantum AG team for their valuable discussions, feedback, and support.

AI Usage Disclosure

Artificial Intelligence tools were used to rephrase and refine the text of this paper. The core software implementation and the original draft of the content were produced by the authors.

References

- Ballester-Ripoll, R. (2022). Tntorch: Tensor network learning with PyTorch. In *GitHub repository*. GitHub. <https://github.com/rballester/tntorch>
- Bohun, V., Lukin, I., Luhanko, M., Korpas, G., De Brouwer, P. J. S., Maksymenko, M., & Koch-Janusz, M. (2024). Entanglement scaling in matrix product state representation of smooth functions and their shallow quantum circuit approximations. *arXiv Preprint arXiv:2412.05202*. <https://arxiv.org/abs/2412.05202>
- Dolgov, S. V., & Savostyanov, D. V. (2014). Alternating minimal energy methods for linear systems in higher dimensions. *SIAM Journal on Scientific Computing*, 36(5), A2248–A2271. <https://doi.org/10.1137/140953289>
- Gribben, D., Allende, C., Villarino, A., Cortines, A., Ali, M., Orús, R., Oswald, P., & Lehdili, N. (2026). STN-GPR: A singularity tensor network framework for efficient option pricing. *arXiv Preprint arXiv:2603.26318*. <https://arxiv.org/abs/2603.26318>

- 153 Guzman, J., & Tiunov, I. (2026). Local interpolation via low-rank tensor trains. *arXiv Preprint*
 154 *arXiv:2601.03885*. <https://arxiv.org/abs/2601.03885>
- 155 Hauschild, J., & Pollmann, F. (2018). Efficient numerical simulations with tensor networks:
 156 Tensor network python (TeNPy). *SciPost Physics Lecture Notes*, 005.
- 157 Ion, I. G., Loukrezis, D., & De Gersem, H. (2022). Tensor train based isogeometric analysis for
 158 PDE approximation on parameter dependent geometries. *Computer Methods in Applied*
 159 *Mechanics and Engineering*, 401, 115593.
- 160 Ion, I. G., & others. (2022). Tensor-train approximation of the chemical master equation and
 161 its application for parameter inference. *Proceedings of MATHMOD 2022*.
- 162 Kossaifi, J., Panagakis, Y., Anandkumar, A., & Pantic, M. (2019). TensorLy: Tensor learning
 163 in python. *Journal of Machine Learning Research*, 20(26), 1–6. [http://jmlr.org/papers/](http://jmlr.org/papers/v20/18-277.html)
 164 [v20/18-277.html](http://jmlr.org/papers/v20/18-277.html)
- 165 Novikov, A., Izmailov, P., Khrulkov, V., Figurnov, M., & Oseledets, I. (2020). Tensor train
 166 decomposition on TensorFlow (T3F). *Journal of Machine Learning Research*, 21(30), 1–7.
- 167 Oseledets, I. (2021). Ttpy: Python implementation of the tensor train (TT)-toolbox. In
 168 *GitHub repository*. GitHub. <https://github.com/oseledets/ttpy>
- 169 Oseledets, I. V. (2010). TT-cross approximation for multidimensional arrays. *Linear Algebra*
 170 *and Its Applications*, 432(1), 70–88. <https://doi.org/10.1016/j.laa.2009.07.024>
- 171 Oseledets, I. V. (2011). Tensor-train decomposition. *SIAM Journal on Scientific Computing*,
 172 33(5), 2295–2317. <https://doi.org/10.1137/090752286>
- 173 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z.,
 174 Gimelshein, N., Antiga, L., Desmaison, A., Kopf, A., Yang, E., DeVito, Z., Raison, M.,
 175 Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., ... Chintala, S. (2019). PyTorch: An
 176 imperative style, high-performance deep learning library. In *Advances in neural information*
 177 *processing systems* 32 (pp. 8024–8035). Curran Associates, Inc. [http://papers.nips.cc/](http://papers.nips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf)
 178 [paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf](http://papers.nips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf)