

macroframe-forecast: Smooth and Constraint-Consistent Forecasting

Sakai Ando¹, Shuvam Das¹, and Sultan Orazbayev¹

¹ International Monetary Fund

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Blake Rayfield](#)

Reviewers:

- [@kaninsin](#)
- [@FATelarico](#)

Submitted: 03 September 2025

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).

Summary

The Python package macroframe-forecast generates forecasts that are both temporally smooth and consistent with user-specified constraints. The framework consists of model estimation and reconciliation: first, unconstrained forecasts are produced using user-specified statistical or machine learning models; second, these forecasts are adjusted to satisfy linear equality and inequality constraints while minimizing volatility over time in the forecast horizon. The package offers a user-friendly interface for specifying constraints as strings and supports advanced customization of weights and smoothing parameters. An example - forecasting U.S. GDP with a terminal GDP growth rate target - illustrates the usage.

Statement of need

In forecasting economic time series, statistical models often need to be supplemented with procedures that impose constraints while preserving smoothness over time. For example, GDP forecasts generated using models such as autoregressions or decision trees may not align with the long-term growth rates anticipated by forecasters. In such cases, forecasters aim to adjust the time series so that it converges smoothly to the desired long-term growth path. However, ad hoc constraint imposition, such as manually altering only the terminal value in a long time series, can introduce undesirable discontinuities between the penultimate and terminal values. Other such examples in economic forecasting can be found in Ando (2024) and Ando & Kim (2023).

Systematically imposing constraints while retaining smoothness is important but challenging. Constraints often stem from accounting identities and expert judgment, making their incorporation essential for internal consistency. Smoothness is equally critical, as optimal forecasts typically exhibit less volatility than historical data, as in random walk. Achieving both objectives manually is resource-intensive, especially when dealing with numerous variables and constraints, raising the question of how to systematically impose constraints and smoothness.

Existing packages in R and Python assist forecast reconciliation and smoothing separately but not jointly. For instance, the hts (Hyndman et al., 2021) and FoReco (Girolimetto & Di Fonzo, 2023) packages in R support reconciliation, but the reconciled forecast may not be smooth over time. This is also the case for hierarchicalforecast (Olivares et al., 2024) package in Python. On the other hand, packages, such as smooth (Svetunkov, 2024) and forecast (Hyndman et al., 2024) for R and statsmodels (Seabold & Perktold, 2010) for Python, provide methods to generate smooth forecasts but do not have the functionality to impose constraints.

To our knowledge, no package supports the simultaneous application of both reconciliation and smoothing, and this is the gap that macroframe-forecast attempts to fill. A more detailed explanation of macroframe-forecast can be found in Ando et al. (2025).

State of the field

Several tools support macroeconomic forecasting, but none let users produce multiple smooth forecasts that also satisfy various constraints. General-purpose econometric packages like Python's statsmodels (Seabold & Perktold, 2010) and R's smooth (Svetunkov, 2024) offer flexible time-series and multivariate models (ARIMA, VAR, VECM), yet enforcing accounting identities or judgmental constraints remains largely manual and ad hoc.

Existing reconciliation packages in R and Python — hts (Hyndman et al., 2021), FoReco (Girolimetto & Di Fonzo, 2023), hierarchicalforecast (Olivares et al., 2024) — can impose constraints, but the resulting reconciled series often display unnatural kinks that are difficult to justify. Enforcing constraints typically disrupts smoothness, while smoothing breaks constraints, meaning both properties must be ensured simultaneously, requiring a dedicated theory and a package that implements it.

macroframe-forecast was developed to fill this gap. It begins with any unconstrained forecast — whether from statistical models in sktime (Löning et al., 2019), expert judgment, or hybrid methods—and systematically enforces high-dimensional accounting identities, inequality constraints, and long-run conditions while preserving temporal smoothness. Users can specify constraints in human-readable strings, which the package automatically converts into the appropriate matrix form, reducing misspecification risk. In this way, macroframe-forecast offers a unique scholarly and practical contribution, complementing rather than competing with existing forecasting or reconciliation tools.

Software Design

The design of macroframe-forecast reflects a trade-off between flexibility, transparency, and computational efficiency. The package adopts a two-step architecture - first generating unconstrained forecasts, then adjusting them through reconciliation and smoothing - which clearly separates model estimation from constraint enforcement. To give users flexibility in choosing forecasting method most suitable for their specific context, we deliberately designed the unconstrained forecasting step to explicitly support any forecasting model implemented in sktime (Löning et al., 2019) without modifying the reconciliation logic.

The second step is formulated as a single quadratic programming problem with an explicit objective function and clearly defined equality and inequality constraints. This approach avoids ad hoc adjustments such as manually applying smoothing after constraints are imposed. Instead, forecast accuracy, smoothness over time, and consistency with constraints are handled jointly within one optimization problem. Constraints are specified using readable string expressions, reducing the need for users to work directly with matrices as in existing packages.

Although existing packages support forecast reconciliation or smoothing individually, there has not been a package that achieves both. Incorporating smoothness penalties into reconciliation-focused packages, or adding general constraint handling to smoothing libraries, would have required substantial structural changes. Building a new package was therefore necessary to implement a transparent, theory-consistent framework that jointly enforces smoothness and constraints while remaining easy to use for applied macroeconomic forecasting.

Research Impact Statement

macroframe-forecast enables economic policy makers to produce multivariate macroeconomic forecasts that are both smooth over time and internally consistent with accounting identities through an intuitive interface. While existing reconciliation tools such as hts (Hyndman et al., 2021) and FoReco (Girolimetto & Di Fonzo, 2023) enforce accounting consistency, they often generate unrealistic kinks in forecast paths, limiting their usefulness in policy settings. Such

artifacts are difficult for policy makers to justify publicly and weaken the credibility of forecasts when communicated to the media.

In the absence of a systematic solution, macroframework forecasting has typically relied on manual adjustments in spreadsheet environments, requiring substantial staff time and limiting scalability and transparency. macroframe-forecast is the first open-source package to jointly ensure smoothness and accounting coherence, allowing experts to incorporate judgment without sacrificing statistical rigor. This significantly lowers the cost of producing explainable, high-dimensional forecasts and improves institutional forecasting capacity.

Community readiness is demonstrated through comprehensive documentation, reproducible example workflows, and a permissive open-source license that encourages reuse and extension. As macroeconomic forecasting increasingly combines expert judgment with data-driven methods, macroframe-forecast provides an infrastructure with strong potential for adoption and downstream impact across economic policy institutions.

AI Usage Disclosure

Generative AI tools (Microsoft Copilot, GitHub Copilot) were used to assist the development of this package, the writing of this manuscript, and the preparation of supporting materials. The majority of the work was produced by the authors, and the usage of AI tools was to facilitate code debugging, generate potential solutions, fix typos and assist with other minor tasks (e.g. formatting). All AI-generated inputs were reviewed and edited by the authors before incorporation in the code/paper.

Method

The framework consists of two steps, where the first step provides users with a flexible choice of forecasting models, and the second step allows users to adjust the first-step forecasts so that the forecasts are smooth over time and satisfy various constraints, such as accounting identities and pre-specified targets. For the theoretical properties, see Ando (2024).

Suppose there are M time series to forecast, and if we stack the forecast horizon of all the time series, the size is N . The first-step forecast $\bar{y} \in \mathbb{R}^N$ can be generated using any model in the `sktime` (Löning et al., 2019) python package. The second-step forecast $\tilde{y}$ is generated by solving the following quadratic programming problem

$$\tilde{y} = \arg \min_{y \in \mathbb{R}^N} (y - \bar{y})^\top W^{-1} (y - \bar{y}) + y^\top \Phi y \quad \text{s.t.} \quad C_{\text{eq}} y - d_{\text{eq}} = 0, \quad C_{\text{ineq}} y - d_{\text{ineq}} \leq 0.$$

The first term penalizes deviations from the first-step forecast $\bar{y}$, weighted by W^{-1} , where W is an estimator of the first step forecast error covariance matrix. It ensures that accurate forecasts are adjusted less than less accurate ones. The second term enforces smoothness via Φ , a block-diagonal matrix built from a degenerate penta-diagonal matrix used in the calculation of the Hodrick-Prescott filter (Hodrick & Prescott, 1997). $C_{\text{eq}} y - d_{\text{eq}} = 0$ refers to equality constraints, and $C_{\text{ineq}} y - d_{\text{ineq}} \leq 0$ refers to inequality constraints.

Weight matrix selection

By default, W is estimated using time series cross validation and Oracle Approximating Shrinkage (OAS) (Chen et al., 2010), which shrinks the sample covariance toward a scalar. An alternative is OAS with diagonal target (Ando & Xiao, 2023), which is robust when variables differ in scale. Users can also specify identity weighting or custom matrices.

127 Smoothness parameters

128 The smoothness matrix Φ is defined as

$$\Phi = \begin{bmatrix} \lambda_1 F_1 & & \\ & \ddots & \\ & & \lambda_M F_M \end{bmatrix}, \quad F_i = \begin{bmatrix} 1 & -2 & 1 & 0 \\ -2 & 5 & -4 & 1 \\ 1 & -4 & 6 & -4 \\ 0 & 1 & -4 & 6 \\ & & \ddots & \ddots & \ddots & \ddots \\ & & & 6 & -4 & 1 & 0 \\ & & & -4 & 6 & -4 & 1 \\ & & & 1 & -4 & 5 & -2 \\ & & & 0 & 1 & -2 & 1 \end{bmatrix}, \quad \lambda_i = \frac{\lambda_i^*}{\sigma_i^2},$$

129 where $i = 1, \dots, M$. Default parameter values λ_i^* follow HP filter conventions (e.g., 1600
130 for quarterly data). It is then scaled by the first step forecast error variance σ_i^2 , obtained
131 from the diagonal elements of the weight matrix W for each time series i , to ensure unit
132 invariance. Setting $\lambda_i^* = 0$ disables smoothing for time series i , reducing the problem to pure
133 reconciliation.

134 An example: single-variable GDP forecast

135 Using U.S. nominal GDP data (1950-2024), we forecast the values for 2025-2030 under the
136 constraint that 2030 growth rate equals 4 percent. When the column name of the variable to
137 be forecasted is GDP, the constraint is expressed as $\text{GDP}_{2030} - 1.04 * \text{GDP}_{2029} = 0$. The
138 MFF class generates the first-step forecast via the default pipeline and reconciles them to satisfy
139 the constraint while smoothing the trajectory. Assuming that the user has read the GDP data
140 as a pandas dataframe named `df0` with index year and GDP, the forecasts (both first and
141 second-step) can be generated as follows:

```
from macroframe_forecast import MFF
from macroframe_forecast.examples import generate_example_GDP_df

df0 = generate_example_GDP_df()
m = MFF(df0, equality_constraints=["GDP_2030 - 1.04 * GDP_2029"])
m.fit()
```




Figure 1: US GDP forecasts: (a) levels and (b) growth.

142 Figure 1: Forecasts of U.S. nominal GDP, 2025-030, under the 4 percent growth constraint in
143 2030.

144 Source: IMF April 2025 World Economic Outlook database and authors' calculations.

145 *Note:* Panel (a) shows the forecast of annual US GDP in trillion USD. Panel (b) shows the
146 growth rates computed from the level forecasts.

147 The second-step forecast can also be customized by providing exogenously defined first-step
148 forecast, weight matrix, or the smoothness parameters. These can be achieved by using the
149 function `Reconciliation` available in the package.

150 Acknowledgements

151 We thank Mingmei Xiao, Yuki Sato, and Doga Bilgin for their earlier contributions.

152 References

- 153 Ando, S. (2024). Smooth Forecast Reconciliation. *IMF Working Papers*, 2024(066), 1.
154 <https://doi.org/10.5089/9798400268922.001>
- 155 Ando, S., Das, S., & Orazbayev, S. (2025). A Python Package to Assist Macroframework
156 Forecasting. *IMF Working Papers*, 2025(172), 1. <https://doi.org/10.5089/9798229023535.001>
- 157
- 158 Ando, S., & Kim, T. (2023). Systematizing macroframework forecasting: High-dimensional
159 conditional forecasting with accounting identities. *IMF Economic Review*, 72(4), 1386–1410.
160 <https://doi.org/10.1057/s41308-023-00225-8>
- 161 Ando, S., & Xiao, M. (2023). High-Dimensional Covariance Matrix Estimation: Shrinkage
162 Toward a Diagonal Target. *IMF Working Papers*, 2023(257), 1. <https://doi.org/10.5089/9798400260780.001>
- 163
- 164 Chen, Y., Wiesel, A., Eldar, Y. C., & Hero, A. O. (2010). Shrinkage algorithms for MMSE
165 covariance estimation. *IEEE Transactions on Signal Processing*, 58(10), 5016–5029.
166 <https://doi.org/10.1109/TSP.2010.2053029>
- 167 Girolimetto, D., & Di Fonzo, T. (2023). *FoReco: Point forecast reconciliation*. <https://cran.r-project.org/package=FoReco>
- 168
- 169 Hodrick, R. J., & Prescott, E. C. (1997). Postwar u.s. Business cycles: An empirical
170 investigation. *Journal of Money, Credit and Banking*, 29(1), 1–16. <http://www.jstor.org/stable/2953682>
- 171
- 172 Hyndman, R., Athanasopoulos, G., Bergmeir, C., Caceres, G., Chhay, L., O'Hara-Wild, M.,
173 Petropoulos, F., Razbash, S., Wang, E., & Yasmeeen, F. (2024). *forecast: Forecasting*
174 *functions for time series and linear models*. <https://pkg.robjhyndman.com/forecast/>
- 175 Hyndman, R., Lee, A., Wang, E., & Wickramasuriya, S. (2021). *Hts: Hierarchical and grouped*
176 *time series*. <https://CRAN.R-project.org/package=hts>
- 177 Lönning, M., Bagnall, A., Ganesh, S., Kazakov, V., Lines, J., & Király, F. J. (2019). *Sktime: A*
178 *unified interface for machine learning with time series*. <https://arxiv.org/abs/1909.07872>
- 179 Olivares, K. G., Garza, A., Luo, D., Challú, C., Mergenthaler, M., Taieb, S. B., Wickramasuriya,
180 S. L., & Dubrawski, A. (2024). *HierarchicalForecast: A reference framework for hierarchical*
181 *forecasting in python*. <https://arxiv.org/abs/2207.03517>
- 182 Seabold, S., & Perktold, J. (2010). Statsmodels: Econometric and statistical modeling with
183 python. *9th Python in Science Conference*.
- 184 Svetunkov, I. (2024). *Smooth: Forecasting using state space models*. <https://CRAN.R-project.org/package=smooth>
- 185