

total-agent-memory: a local, reproducible long-term memory server for AI coding agents

Vitalii Cherepanov ¹

¹ Independent researcher, Novi Sad, Serbia

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: 

Submitted: 28 September 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

AI coding agents such as Claude Code, Codex CLI and Cursor start every session without memory of earlier ones. Decisions, fixes and project conventions from yesterday have to be explained again. total-agent-memory (TAM) is an open-source memory server for such agents. It stores decisions, solutions, facts, errors and session summaries on the user's machine and gives them back through the Model Context Protocol (MCP) ([Model Context Protocol, 2026](#)), so any MCP client can use it without code changes.

Each TAM store is one directory built around a SQLite database. Search combines full-text BM25 ([Robertson & Zaragoza, 2009](#)), dense embeddings computed locally, fuzzy matching and a knowledge graph, and fuses the ranked lists with reciprocal rank fusion ([Cormack et al., 2009](#)). An optional cross-encoder reranks the result and maximal marginal relevance ([Carbonell & Goldstein, 1998](#)) can diversify it. The default profile makes no LLM call on write or search. Facts can carry validity intervals, so the store can answer what was true at a given time, and a newer value of a single-valued fact can retire the older one. Since release 14.6.0 the same code can also run as a team server with separate personal, department and company areas.

The package is written in Python, released under the MIT license, and distributed on PyPI as total-agent-memory.

Statement of need

Research on agent memory needs systems that can be run, inspected and measured on public benchmarks such as LoCoMo ([Maharana et al., 2024](#)), LongMemEval ([Wu et al., 2024](#)), BEAM ([Tavakoli et al., 2025](#)) and MemoryAgentBench ([Hu et al., 2025](#)). Many memory layers that are used in practice are hosted services, or they call an LLM on every write to extract memories ([Chhikara et al., 2025](#)). That makes experiments slower and costlier, ties results to a provider's model versions, and makes runs hard to repeat.

TAM is built for this setting. Its default pipeline is deterministic and runs offline, so retrieval quality can be measured without any model in the loop and a rerun gives the same result. Every runner for the public benchmarks disables usage-based score boosts, which would otherwise let a store score better on its own second run. Negative controls (random, earliest and most recent turns) are scored next to the main metric on LoCoMo, so each reported number comes with its floor. When an LLM is used, for example to answer questions in end-to-end evaluations, the model, prompt and seed are recorded and several runs are reported.

The target users are researchers and engineers who study memory for agents, compare retrieval designs, or need a memory baseline that they can run locally and change. The same server is used in daily work with coding agents, which brings real usage and bug reports.

State of the field

MemGPT (Packer et al., 2023), now continued as Letta, manages tiers of context for one agent. Mem0 (Chhikara et al., 2025) extracts and consolidates facts from conversations with an LLM, with an optional graph variant. Zep (Rasmussen et al., 2025) builds a temporal knowledge graph of conversations and business data. These systems focus on what to store and how to retrieve it, and they are usually deployed as services with an LLM in the write path.

TAM differs in three ways. First, it is local by default and needs no LLM to write or search, which is what makes deterministic retrieval benchmarks possible. Second, it is MCP-native: the memory is a tool server that any MCP client can call, rather than a library that has to be wired into one agent framework. Third, it treats evaluation as part of the software. The repository ships runners and raw results for LoCoMo, LongMemEval, BEAM and MemoryAgentBench, a head-to-head report that grades TAM's answers and Mem0 Platform's published answers under the same judge, and a benchmark of department isolation on the team server (Cherepanov, 2026).

We do not claim that TAM retrieves better than the systems above in general. The head-to-head report found no statistically detected difference from Mem0 Platform on the held-out questions it covers, and it states that this does not show equivalence.

Example

An agent uses TAM through ordinary MCP tool calls. In one session the user explains a choice, and the agent saves it as a decision with the reason:

```
memory_save(type="decision",
            content="Chose pgvector over ChromaDB for the vector store",
            context="WHY: one Postgres instance, row-level security per tenant",
            project="billing")
```

Days later, in a new session and possibly another directory, the user asks why pgvector was chosen. The agent calls:

```
memory_recall(query="vector database choice", project="billing")
```

and gets the decision back with its reason, the time it was saved and its ID. If the decision changes, `memory_update` stores a new version, the next recall returns it, and `memory_history` still shows the old one. The benchmark runners in the repository use the same calls, so experiments exercise the code path that users run.

65 Software design




Figure 1: Main parts of total-agent-memory. Agents call MCP tools; writes and searches go to one store per memory area, built around a SQLite database.

66 **Figure 1** shows the main parts. The design follows from three goals: memory must stay on the
67 user's machine, results must be repeatable, and any MCP client must be able to use it.

68 **One store, one directory.** A store is a directory built around a SQLite database with FTS5
69 for text search and tables for records, versions and the knowledge graph. Embeddings are
70 computed with FastEmbed on the CPU and stored in SQLite next to the records, both as full
71 float32 vectors and as binary-quantized copies for fast candidate search. A ChromaDB index
72 in the same directory serves as a fallback while the SQLite vector index is not yet populated. A
73 store in one directory is easy to copy, back up, delete and reset between experiments. On the
74 team server, a change, its new version and its audit row are written in one SQLite transaction.
75 The cost is one writer per database at a time, which is enough for one person and is handled
76 by one worker process per area on the team server.

77 **No LLM in the default path.** TAM stores what the agent sends and indexes it locally instead
78 of asking an LLM to extract facts. Writes stay fast and free, data stays on the machine,
79 and retrieval is deterministic, so a benchmark run can be repeated exactly. An LLM can be
80 configured for optional stages such as summaries, contradiction checks or query rewriting, and
81 each of them can be switched off.

82 **Hybrid retrieval with fusion.** Reciprocal rank fusion needs no score calibration between the
83 BM25, dense, fuzzy and graph retrievers. The fused pipeline answers in tens of milliseconds
84 on a warm store; a cold first query pays for loading the models.

85 **Facts that change.** Records are versioned instead of overwritten, and knowledge-graph facts
86 carry valid_from and valid_to, so a query can ask what was true at a point in time. The

87 store stays append-only and auditable, at the cost of a larger database.

88 **MCP as the interface.** The server exposes its operations as MCP tools over stdio or HTTP.
89 An index mode of `memory_recall` returns only IDs and titles, and `memory_get` fetches the few
90 records the agent needs, which keeps its context small.

91 **Team server.** For companies, a gateway authenticates personal tokens and runs each personal,
92 department and shared area in its own worker process and store, checking membership on
93 every request. On a synthetic company, attacks across nine attempt types returned no records
94 of other departments, and membership changes applied from the next request (Cherepanov,
95 2026). The cost is one process per active area; release 14.6.0 moved the cross-encoder to the
96 gateway so that it runs once per query instead of once per area. A PostgreSQL backend with
97 one schema and database role per area is also available.

98 **Tests.** The repository has 255 test files. Continuous integration installs and smoke-tests the
99 package on Linux, macOS and Windows, and runs the team-server suite on both SQLite and
100 PostgreSQL.

101 Research impact statement

102 TAM is used as a method in MemoryAgentBench (Hu et al., 2025): its maintainers merged
103 an adapter that runs TAM under the same configurations as the other memory methods in
104 the benchmark (HUST-AI-HYZ/MemoryAgentBench pull request 25, September 2026). An
105 independent atlas of open-source agent-memory systems described one TAM technique, a
106 second retrieval that looks for contradicting evidence before answering, as worth reusing in
107 other systems. The saved verdicts of the head-to-head report were recomputed by an external
108 reviewer from the published files, and the report was revised after that review; the revision
109 history is in the report.

110 The repository has received pull requests from outside contributors, and the software has been
111 forked by other developers. The team-server isolation study built on TAM is published as a
112 preprint with its data and harness (Cherepanov, 2026).

113 AI usage disclosure

114 Generative AI tools were used throughout this project. AI coding assistants (Claude Code
115 with Anthropic Claude models, and OpenAI Codex) were used to write and refactor code,
116 write tests and benchmark harnesses, and draft documentation. An AI assistant (Claude Code
117 with Claude Opus 5.5) helped draft this paper and check its statements against the repository.
118 The author designed the system and chose its architecture, reviewed the generated code and
119 text, decided what to measure and report, and is responsible for the accuracy of all submitted
120 material.

121 Acknowledgements

122 I thank the outside contributors who sent pull requests and bug reports, and the reviewer who
123 recomputed the head-to-head verdicts.

124 References

125 Carbonell, J., & Goldstein, J. (1998). The use of MMR, diversity-based reranking for reordering
126 documents and producing summaries. *Proceedings of the 21st Annual International*
127 *ACM SIGIR Conference on Research and Development in Information Retrieval*, 335–336.
128 <https://doi.org/10.1145/290941.291025>

- Cherepanov, V. (2026). *Organizational memory for AI agents: Department scopes, onboarding and offboarding in a multi-tenant memory server*. Zenodo. <https://doi.org/10.5281/zenodo.23011523>
- Chhikara, P., Khant, D., Aryan, S., Singh, T., & Yadav, D. (2025). *Mem0: Building production-ready AI agents with scalable long-term memory*. <https://doi.org/10.48550/arXiv.2504.19413>
- Cormack, G. V., Clarke, C. L. A., & Büttcher, S. (2009). Reciprocal rank fusion outperforms Condorcet and individual rank learning methods. *Proceedings of the 32nd International ACM SIGIR Conference on Research and Development in Information Retrieval*, 758–759. <https://doi.org/10.1145/1571941.1572114>
- Hu, Y., Wang, Y., & McAuley, J. (2025). *Evaluating memory in LLM agents via incremental multi-turn interactions*. <https://doi.org/10.48550/arXiv.2507.05257>
- Maharana, A., Lee, D.-H., Tulyakov, S., Bansal, M., Barbieri, F., & Fang, Y. (2024). *Evaluating very long-term conversational memory of LLM agents*. <https://doi.org/10.48550/arXiv.2402.17753>
- Model Context Protocol. (2026). *Model context protocol specification*. <https://modelcontextprotocol.io/specification>
- Packer, C., Wooders, S., Lin, K., Fang, V., Patil, S. G., Stoica, I., & Gonzalez, J. E. (2023). *MemGPT: Towards LLMs as operating systems*. <https://doi.org/10.48550/arXiv.2310.08560>
- Rasmussen, P., Paliychuk, P., Beauvais, T., Ryan, J., & Chalef, D. (2025). *Zep: A temporal knowledge graph architecture for agent memory*. <https://doi.org/10.48550/arXiv.2501.13956>
- Robertson, S., & Zaragoza, H. (2009). *The probabilistic relevance framework: BM25 and beyond* (Vol. 3, pp. 333–389). Now Publishers. <https://doi.org/10.1561/15000000019>
- Tavakoli, M., Salemi, A., Ye, C., Abdalla, M., Zamani, H., & Mitchell, J. R. (2025). *Beyond a million tokens: Benchmarking and enhancing long-term memory in LLMs*. <https://doi.org/10.48550/arXiv.2510.27246>
- Wu, D., Wang, H., Yu, W., Zhang, Y., Chang, K.-W., & Yu, D. (2024). *LongMemEval: Benchmarking chat assistants on long-term interactive memory*. <https://doi.org/10.48550/arXiv.2410.10813>